

Triage Verdict

ALERT ID: CLD-0101 **SOURCE SYSTEM:** Identity Monitoring Entra ID **ALERT TIME (UTC):** 2026-08-31 23:41:07 **ANALYST:** Uju Igweokwu

ALERT:

18 sing-in failures on svc-backup@cloudora.io from 203.0.113.20 (LDN-SRV-BK1), 31 Aug 23:41 to 1 Sep 00:03 followed by a successful sign-in at 00:14 UTC Repeated authentication reported high.

HYPOTHESIS:

Malicious: an internal brute-force attack against a service account.

Benign: Repeated authentication failure is from an announced scheduled backup credential rotation

EVIDENCE CHECKED:

The shift handover document change record CHG-2088 stated a scheduled backup credential rotation at 23:00 to 02:30 and warns about expected failure, the sign-in record shows 18 failed authentication error 50126 from LDN-SRV-BK1 The backup server's own address.

VERDICT

False
positive

SEVERITY YOU ASSIGNED

Informational

SEVERITY:

Reported High but was rerated as Informational because the activity was a scheduled approved maintenance window and the access only involved the backup server.

ACTION:

closed with reference to CHG-2088, and recommends suppressing service-account failures during approved change windows.

JUSTIFICATION:

Source host, account, time window and recovery all match the announced change; brute force does not stop at 18 attempts, does not originate from the target's own server, and does not resolve exactly when a rotation job completes.

MITRE ATT&CK:

None. False positive gets no MITRE technique.

Triage Verdict

ALERT ID: CLD-0102 **SOURCE SYSTEM:** Endpoint Protection (Defender AV) **ALERT TIME (UTC):** 2026-09-01 02:45:33 **ANALYST:** Uju Igweokwu

ALERT:

Trojan:Win32/Stealgen.B on LDN-WS-117 (liam.doyle), four detections of malware at 02:13 to 02:41 UTC, quarantine failed each time. Reported High.

HYPOTHESIS:

Malicious: An active credential-stealing malware resisting removal.

Benign: false or stale detection on an already-clean host.

EVIDENCE CHECKED:

AV log: scheduled scan clean on 30 Aug; four detections on 1 Sep with escalating failure reasons (file in use, access denied, file recreated) and a second dropped file wsync_upd.exe at 02:41 in the same AppData folder. Handover: no change or deployment explains it. Sign-in slice: no anomalous sign-ins for liam.doyle yet (normal 08:05 logon).

VERDICT

True positive

SEVERITY YOU ASSIGNED

Sev 2

SEVERITY:

Reported High, credential stealing malware infection confirmed and was assigned Sev 2

ACTION:

Isolate LDN-WS-117 from the network via endpoint management,
Raise incident for eradication/reimage,
Force password reset and session revocation for liam.doyle as a precaution,
Monitor for anomalous sign-ins.

JUSTIFICATION:

Repeated failed quarantine and file regeneration under another name confirms an active malware infection A false detection does not fail quarantine three different ways and then recreate itself under a new name.

MITRE ATT&CK:

T1204.002 (User Execution: Malicious File).

Triage Verdict

ALERT ID: CLD-0103 **SOURCE SYSTEM:** IDS **ALERT TIME (UTC):** 2026-09-01 02:52:010 **ANALYST:** Uju Igweekwu

ALERT:

IDS:203.0.223.44 Probed over 4,000 TCP port across 180 internal host.

HYPOTHESIS:

Malicious: A compromised host performing unauthorised internal reconnaissance.

Benign: An authorised scheduled vulnerability scan.

EVIDENCE CHECKED:

Handover schedule: LDN-SCAN-01 (203.0.113.44) sweeps Monday 22:00 to Tuesday 03:00; source and window match exactly. Alert detail: sequential SYN probes, no data transfer, no follow-on alerts from the source.

VERDICT

False
positive

SEVERITY YOU ASSIGNED

Informational

SEVERITY:

Reported Medium but was rerated as Informational because the activity was a scheduled approved scan.

ACTION:

Closed with reference to the authorised scan schedule.

Recommend a tuning change via change control to suppress the port-sweep rule for LDN-SCAN-01 during its approved window; and not to disable the rule globally.

JUSTIFICATION:

Source, scope, timing and traffic shape all match the documented weekly scan; attacker recon pivots onto hosts of interest and moves data, which is absent here.

MITRE ATT&CK:

None. False positive gets no MITRE technique.

Triage Verdict

ALERT ID: CLD-0104 **SOURCE SYSTEM:** Sign-in logs (Entra ID) **ALERT TIME (UTC):** 2026-08-13 06:31
ANALYST: Uju Igweokwu

ALERT:

Five MFA denials 06:31 to 06:47 UTC, success 06:52, new authenticator registered 06:50. Reported Medium.

HYPOTHESIS:

Malicious: Attempted MFA bypass attack followed by a successful sign-in.

Benign: User error, failing MFA before successful authentication on new phone.

EVIDENCE CHECKED:

Sign-in log: every event, failures included, from her usual Manchester office IP on her usual Windows 10 Chrome machine, in her normal morning window. Helpdesk HD-4471: raised by nina.cole herself at 07:05 asking for help re-enrolling MFA on a new phone. Registration timing (06:50) sits between her failures and her success, consistent with authentication finally working.

VERDICT

False
positive

SEVERITY YOU ASSIGNED

Informational

SEVERITY:

Reported Medium but was rerated as Informational because the activity was consistent with the user's normal behaviour.

ACTION:

No password reset required. Close referencing HD-4471. Recommend auto-annotating MFA re-enrolment events when a matching helpdesk ticket exists.

JUSTIFICATION:

The event matches the user's baseline behaviour.

MITRE ATT&CK:

None. False positive gets no MITRE technique.

Triage Verdict

ALERT ID: CLD-0105 **SOURCE SYSTEM:** User report (Email Security) **ALERT TIME (UTC):** 2026-08-31 23:41:07 **ANALYST:** Uju Igweokwu

ALERT:

User report from james.holt: payroll-branded email with subject "Your August payslip is ready" and a portal link.

HYPOTHESIS:

Malicious: A second Phishing email attempt.

Benign: A legitimate monthly payroll notification from Cloudora's HR

EVIDENCE CHECKED:

Raw email headers show the from domain is cloudora.io with SPF, DKIM and DMARC all passing and aligned to cloudora.io.

The Received chain shows the message originated internally from myhr-notify.cloudora.io (203.0.113.8) and did not leave Cloudora infrastructure.

Email body links only to the legitimate myhr.cloudora.io portal and contains no credential, payment or other suspicious request

VERDICT

False
positive

SEVERITY YOU ASSIGNED

Informational

SEVERITY:

Reported as suspicious but was rerated as Informational because the message was confirmed as legitimate.

ACTION:

closed the alert, replied to the reporter thanking him for reporting it, confirm the email is genuine, and encourage continued reporting.

JUSTIFICATION:

Email headers showed the from domain is cloudora.io with SPF, DKIM and DMARC all passed and aligned to cloudora.io

MITRE ATT&CK:

None. False positive gets no MITRE technique.

Triage Verdict

ALERT ID: CLD-0106 **SOURCE SYSTEM:** Endpoint Protection Dashboard **ALERT TIME (UTC):** 08:00
ANALYST: Uju Igweokwu

ALERT:

08:05 UTC: one device with active unresolved threats (LDN-WS-117, Stealgen.B).

HYPOTHESIS:

Malicious: A second endpoint compromise.

Benign: A duplicate alert of CLD-0102

EVIDENCE CHECKED:

Defender AV daily digest at 08:00 summarises the same LDN-WS-117 Trojan: Win32/ Stealgen.B detections already handled under CLD-0102. No new host, detection, indicator, or time window identified.

VERDICT

Duplicate

SEVERITY YOU ASSIGNED

Informational

SEVERITY:

As in CLD-0102

ACTION:

closed as duplicate and cross-reference CLD-0102. Recommend correlating the Defender digest connector with real-time alerts to prevent duplicate tickets and preserve a single evidence trail.

JUSTIFICATION:

Event must have one ticket carrying the evidence trail; the digest adds a delivery path, not information.

MITRE ATT&CK:

As in CLD-0102

Triage Verdict

ALERT ID: CLD-0107 **SOURCE SYSTEM:** SherePoint Monitoring **ALERT TIME (UTC):** 2026-09-01 03:00
ANALYST: Uju Igweokwu

ALERT:

The rule fired because it detected a large number of deletions 1,240 files were deleted from HR-Records-Archive, folder Leavers-2019, 03:00 to 03:20 UTC under tara.kemp's account

HYPOTHESIS:

Malicious: A compromised account deleting a large number of HR records

Benign: A legitimate HR data deletion.

EVIDENCE CHECKED:

The handover confirms HR-Retention Flow is executing this week under Tara Kemp's account during the early hours to delete expired records from the HR archive under the seven-year retention policy. The 1,240 deletions occurred in HR-Records-Archive at 03:00 and match the expected retention activity and initiating process.

VERDICT

False
positive

SEVERITY

Informational

SEVERITY:

Reported as High re-rated low as Informational, authorised automation, not compromised-account or ransomware activity.

ACTION:

Close with reference to the retention schedule and recommend tuning the mass-deletion rule to recognise HR-RetentionFlow operating on HR-Records-Archive during its approved retention window, while continuing to alert on mass deletion outside that process, folder or window.

JUSTIFICATION:

The large file deletion was initiated by the scheduled HR-RetentionFlow automation.

MITRE ATT&CK:

None. False positive gets no MITRE technique.

Triage Verdict

ALERT ID: CLD-0108 **SOURCE SYSTEM:** Identity Monitoring **ALERT TIME (UTC):** 03:54 **ANALYST:** Uju Igweokwu

ALERT:

First-time sign-in properties for gwen.muir at 03:54 UTC, from 198.18.163.41 (London, residential provider), client IMAP4 (legacy), Windows 10, Firefox. No preceding failures.

HYPOTHESIS:

Malicious: A compromised credentials.

Benign: A legitimate sign-in from a new device.

EVIDENCE CHECKED:

sign-in slice baseline for gwen.muir, including normal device, browser, applications and hours; alert event at 03:54 from a London residential IP using Windows 10, Firefox and legacy IMAP4.

VERDICT

True positive

SEVERITY

Sev 2

SEVERITY:

Reported as Low re-rated upward the successful first-time legacy-protocol sign-in is suspicious and may indicate already compromised credentials.

ACTION:

Immediately disable or block the account, revoke active sessions/tokens, reset the password, disable legacy IMAP4 authentication for the account, and investigate for additional account activity. Account escalated to the Identity and Access Management team to gather authentication logs, recent password-reset/change history, mailbox access/audit logs, device and application sign-in details, and other recent sign-ins to determine how the credential was obtained and whether the account was accessed elsewhere.

JUSTIFICATION:

The event differs from Gwen's established device, browser, application and time-of-day baseline despite the London location matching. The successful first attempt with no preceding failures is consistent with an already-obtained credential, while legacy IMAP4 represents a policy gap because it may bypass modern MFA controls.

MITRE ATT&CK:

T1078.004 (Valid Accounts: Cloud Accounts), via legacy protocol to bypass MFA.

Triage Verdict

ALERT ID: CLD-0109 **SOURCE SYSTEM:** Network Security Monitoring
ANALYST: Uju Igweokwu

ALERT TIME (UTC): 07:40

ALERT:

Outbound connections to uncategorised destination with a domain registered 9 days ago, LDN-WS-104 (aria.reid), MAN-LT-033 (finn.gale), AUS-WS-208 (dean.page): outbound HTTPS to 198.18.207.66 every 60 seconds since 07:40 UTC.

HYPOTHESIS:

Malicious: A Potential malicious beaconing across three offices.

Benign: An undocumented agent or service.

EVIDENCE CHECKED:

LDN-WS-104, MAN-LT-033 and AUS-WS-208 each connected to the same nine-day-old external host every 60 seconds since 07:40, indicating a possible beaconing pattern. Sign-in activity for the three assigned users showed no anomalous authentication activity.

Handover: no change record or rollout explains a new mass-deployed agent.

VERDICT

Escalate

SEVERITY

Sev 3

SEVERITY:

Reported as Medium re-rated High for further investigation; activity cannot be safely closed or resolved at triage.

ACTION:

Escalate immediately to the vCISO with the three affected hostnames, destination host, connection interval and start time, relevant network/endpoint telemetry, and the sign-in evidence showing no anomalous user activity. Request investigation of the destination, process and persistence on each host and correlation across the three sites. Containment of the three hosts is deferred pending the appropriate incident-response decision as containment decisions affecting multiple hosts across three sites require the responder/vCISO decision.

JUSTIFICATION:

The pattern of three hosts in three offices communicating with the same rare external host at exactly 60-second intervals is consistent with beaconing/C2-style activity. The fact that the sign-in slice does not show anomalous user activity, the strongest evidence is in the network/endpoint telemetry, not the identity logs.

MITRE ATT&CK:

Suspected T1071.001 (Application Layer Protocol: Web Protocols), pending confirmation.

Triage Verdict

ALERT ID: CLD-0110 **SOURCE SYSTEM:** Watchlist IP Activity
ANALYST: Uju Igweokwu

ALERT TIME (UTC): 05:09- 05:11

ALERT:

The source IP was already present on the tenant block list and had previously been identified resulting in the three sign-in attempts for daniel.reeve from 102.89.44.17 (Lagos; CLD-0001 attacker infrastructure), 05:09 to 05:11 UTC, every one returning 53003, blocked by Conditional Access.

HYPOTHESIS:

Malicious: A Known threat actor attempting to authenticate.

Benign: Security controls successfully prevented access of a blocked credential.

EVIDENCE CHECKED:

three sign-in attempts for daniel.reeve from 102.89.44.17, a known CLD-0001 attacker infrastructure IP, all returned error 53003 and were blocked by Conditional Access. No successful sign-in from 102.89.x was identified in the slice, and the CEO's normal morning activity remained unaffected. The attempted activity and confirmed block to the CLD-0001 record as evidence of infrastructure reuse.

VERDICT

**True
Positive**

SEVERITY

Sev 1

SEVERITY:

Reported as Low retained because the control prevented successful access.

ACTION:

Continue monitoring the watchlist IP and re-raise severity if any sign-in from 102.89.x succeeds or Conditional Access no longer blocks activity

JUSTIFICATION:

53003 proves the requests died at policy; there is no evidence of impact anywhere in the data. But attacker infrastructure re-use against the same executive is intelligence: it costs nothing to record and it changes how fast the next Lagos alert gets treated.

MITRE ATT&CK:

T1110.003 (Password Spraying), attempted

Triage Verdict

ALERT ID: CLD-0111 **SOURCE SYSTEM:** Reported Email
lgweekwu

ALERT TIME (UTC): 08:19 **ANALYST:** Uju

ALERT:

At 08:19 UTC lucas.ford@cloudora.io (accounts payable) reports an invoice reminder from billing@meridianoffice.example the subject was "Overdue invoice INV-2214, second reminder." And a PDF attachment; he is not sure Cloudora uses the supplier, and does not think he opened the attachment.

HYPOTHESIS:

Malicious: A phishing email.

Benign: A legitimate second reminder of an overdue invoice.

EVIDENCE CHECKED:

Email headers show SPF pass, but DKIM and DMARC are absent, which does not establish that the message is either legitimate or malicious. The sender/domain, supplier relationship, attachment contents, recipient scope and whether the attachment was opened cannot currently be confirmed from the available evidence.

VERDICT

**Insufficient
data**

SEVERITY

Sev 3

SEVERITY:

Reported as Medium re-rated to High neither phishing nor legitimate mail is defensible at this stage.

ACTION:

Verify Meridian Office Supplies with Accounts Payable/procurement, safely detonate or scan the PDF in the approved malware-analysis/sandbox tool, search mail logs/message trace to identify other recipients, and check endpoint/application telemetry to determine whether Lucas Ford opened or executed the attachment.

JUSTIFICATION:

Reported supplier invoice requires further verification, the user's interaction is unverified.

MITRE ATT&CK:

None yet; T1566.001 (Spearphishing Attachment) if confirmed True Positive.

Triage Verdict

ALERT ID: CLD-0112 **SOURCE SYSTEM:** Identity Monitoring Entra ID **ALERT TIME (UTC):** 05:58 **ANALYST:** Uju Igweokwu

ALERT:

Sign-in from Dubai 05:58 UTC, then London (203.0.113.10) 06:40 UTC, 42 minutes apart fired alert impossible travel.

HYPOTHESIS:

Malicious: An attacker signing-in with stolen credential.

Benign: A legitimate travel with non-interactive office desktop activity.

EVIDENCE CHECKED:

Omar Farah's sign-in slice shows the Dubai event at 05:58 as the genuine interactive user activity while the London event at 06:40 is a non-interactive token refresh from his docked office desktop. The London event therefore represents a machine at a fixed desk refreshing a token rather than Omar physically signing in from London. Omar's documented Dubai travel history in CLD-0001 supports the travel context.

VERDICT

False
Positive

SEVERITY

Informational

SEVERITY:

Reported as Medium re-rated to Informational legitimate travel with expected non-interactive desktop activity.

ACTION:

Closed with reference to the documented travel context and recommend a low-cost travel-notification hygiene improvement so expected travel can be identified during triage.

JUSTIFICATION:

The misleading event is the London non-interactive token refresh, while the Dubai event is the interactive human activity. A true positive would have been indicated by an interactive London sign-in while Omar was in Dubai or a different device/application pattern inconsistent with his known activity.

MITRE ATT&CK:

None. False positive gets no MITRE technique.

Task 2 model: The End of Shift Summary

Alert summary

Alert	Verdict	Sev	Action
CLD-0101	False positive maintenance	Informational	Closed; tuning recommended. Resolved
CLD-0102	True positive: active credential stealer	Sev 2	Isolated; incident raised; user creds reset. In Progress
CLD-0103	False positive (authorised scanner)	Informational	Closed; scoped IDS exception recommended. Resolved
CLD-0104	False positive MFA phone replacement	Informational	Closed. Resolved
CLD-0105	False positive legitimate payroll mail	Informational	Closed; thanked reporter. Resolved
CLD-0106	Duplicate of CLD-0102	Informational	Closed; connector correlation recommended. Resolved
CLD-0107	False positive (HR retention flow)	Informational	Closed; recommend narrow retention-rule tuning. Resolved
CLD-0108	True positive: credential compromise	Sev 2	In Progress; incident; contain account and investigate
CLD-0109	Escalated (multi-site beaconing)	Sev 3	Escalate to vCISO; containment deferred. In Preogress
CLD-0110	TP activity, blocked; monitoring	Sev 1	Closed; monitoring watchlist infrastructure. Resolved
CLD-0111	Insufficient data; gathering	Sev 3	On hold while awaiting evidence
CLD-0112	False positive (travel + token refresh)	Informational	Closed; recommended travel notification hygiene. Resolved

Task 3 model: The Tuning Table

Rule	Problem	Proposed change	Risk of the change
HR mass deletion	HR-RetentionFlow generates legitimate mass-deletion alerts every retention cycle	Add a narrow exception for HR-RetentionFlow + HR-Records-Archive + approved retention window	A badly scoped exception could hide genuine deletion activity; keep folder, process and time conditions together
Internal port-scan rule	LDN-SCAN-01's authorised weekly vulnerability sweep repeatedly triggers the port-scan rule	Suppress/annotate only LDN-SCAN-01 during its approved scan window	A compromised scanner could evade detection if the exception is too broad; never disable the rule globally
Defender daily-digest connector	The daily digest creates a second ticket for an already-alerted active threat	Correlate digest events against existing real-time incidents, or suppress duplicates	A faulty correlation could suppress a genuinely new detection
MFA re-enrolment alerts	Legitimate new-phone MFA re-enrolment creates attack-shaped alerts	Auto-annotate/raise confidence when a matching helpdesk ticket exists	A fake/helpdesk-compromised ticket could provide false reassurance
Impossible-travel rule	Non-interactive token refreshes from fixed office desktops create misleading geographic events	Distinguish interactive human sign-ins from non-interactive token refreshes	Over-filtering could hide genuine token/session abuse
Travel-notification process	Legitimate travel creates avoidable impossible-travel investigations	Encourage users to notify the SOC/helpdesk of planned travel	Users may forget to notify; notification should reduce noise, not suppress alerts
Watchlist correlation	Known attacker infrastructure continues to generate individual alerts	Correlate watchlist activity with prior incidents such as CLD-0001 and automatically surface the history	Over-correlation could obscure a new campaign if analysts assume every hit is merely historical

Task 4 model: Answer

Which alert was the most dangerous and why?

CLD-0108 was the most dangerous because it represents a probable credential compromise that initially arrived with a misleading Low severity. The successful authentication used legacy IMAP4, differed from Gwen's normal baseline and occurred without preceding failures, meaning an attacker may already possess the password and potentially bypass modern MFA protections.

Which alert would you have got wrong a month ago, and what saved you today?

CLD-0112 is the alert I might have incorrectly called a true positive because "Dubai to London in 42 minutes" looks immediately malicious. What saved me was slowing down and checking the event type, device and application: the Dubai event was interactive human activity while the London event was a non-interactive token refresh from Omar's docked office desktop.